

LAB RECORD

1BM19CS194

Anitej Prasad

6-D

LAB7

```
hduser@bmsce-Precision-T1700:~$ start-all.sh
```

This script is Deprecated. Instead use start-dfs.sh and start-yarn.sh

Starting namenodes on [localhost]

```
hduser@localhost's password:
```

localhost: namenode running as process 10473. Stop it first.

```
hduser@localhost's password:
```

localhost: datanode running as process 10644. Stop it first.

Starting secondary namenodes [0.0.0.0]

```
hduser@0.0.0.0's password:
```

0.0.0.0: secondarynamenode running as process 10857. Stop it first.

starting yarn daemons

resourcemanager running as process 9796. Stop it first.

```
hduser@localhost's password:
```

localhost: nodemanager running as process 10160. Stop it first.

```
hduser@bmsce-Precision-T1700:~$ jps
```

10160 NodeManager

7441 org.eclipse.equinox.launcher_1.5.600.v20191014-2022.jar

9796 ResourceManager

12692 org.eclipse.equinox.launcher_1.5.600.v20191014-2022.jar

10644 DataNode

10857 SecondaryNameNode

10473 NameNode

15100 Jps

```
hduser@bmsce-Precision-T1700:~$ hadoop fs -ls /
```

Found 10 items

```
drwxr-xr-x - hduser supergroup      0 2019-10-23 09:52 /sample
drwxr-xr-x - hduser supergroup      0 2019-10-23 10:33 /test
drwxr-xr-x - hduser supergroup      0 2022-06-14 10:50 /tmp1
drwxr-xr-x - hduser supergroup      0 2019-10-23 09:58 /output
drwxr-xr-x - hduser supergroup      0 2022-06-15 10:27 /rgs
drwxr-xr-x - hduser supergroup      0 2019-10-23 11:09 /stud
drwxr-xr-x - hduser supergroup      0 2019-10-23 15:50 /testing
drwxrwxr-x - hduser supergroup      0 2019-10-23 11:24 /tmp
drwxr-xr-x - hduser supergroup      0 2019-08-01 16:03 /user
```

```
hduser@bmsce-Precision-T1700:~$ hadoop fs -mkdir /1BM19CS194
```

```
hduser@bmsce-Precision-T1700:~$ hadoop fs -copyFromLocal
/home/hduser/Desktop/sample1.txt /1BM19CS194/test.txt
```

```
hduser@bmsce-Precision-T1700:~$ hdfs dfs -cat /1BM19CS194/test.txt
```

hi my name is anitej

my age is twenty one

my country is india

my surname is prasad

```
hduser@bmsce-Precision-T1700:~$ hadoop jar  
/home/hduser/Documents/wordCount.jar wordCount.WCDriver  
/1BM19CS194/test.txt /1BM19CS194/output
```

22/06/15 10:27:53 INFO Configuration.deprecation: session.id is deprecated.
Instead, use dfs.metrics.session-id

22/06/15 10:27:53 INFO jvm.JvmMetrics: Initializing JVM Metrics with
processName=JobTracker, sessionId=

22/06/15 10:27:53 INFO jvm.JvmMetrics: Cannot initialize JVM Metrics with
processName=JobTracker, sessionId= - already initialized

22/06/15 10:27:53 WARN mapreduce.JobSubmitter: Hadoop command-line
option parsing not performed. Implement the Tool interface and execute your
application with ToolRunner to remedy this.

22/06/15 10:27:53 INFO mapred.FileInputFormat: Total input paths to process
: 1

22/06/15 10:27:53 INFO mapreduce.JobSubmitter: number of splits:1

22/06/15 10:27:53 INFO mapreduce.JobSubmitter: Submitting tokens for job:
job_local1115189753_0001

22/06/15 10:27:53 INFO mapreduce.Job: The url to track the job:
<http://localhost:8080/>

22/06/15 10:27:53 INFO mapred.LocalJobRunner: OutputCommitter set in
config null

22/06/15 10:27:53 INFO mapreduce.Job: Running job:
job_local1115189753_0001

22/06/15 10:27:53 INFO mapred.LocalJobRunner: OutputCommitter is
org.apache.hadoop.mapred.FileOutputCommitter

22/06/15 10:27:53 INFO mapred.LocalJobRunner: Waiting for map tasks

22/06/15 10:27:53 INFO mapred.LocalJobRunner: Starting task:
attempt_local1115189753_0001_m_000000_0

22/06/15 10:27:53 INFO mapred.Task: Using ResourceCalculatorProcessTree :
[]

22/06/15 10:27:53 INFO mapred.MapTask: Processing split:
hdfs://localhost:54310/rgs/test.txt:0+89

22/06/15 10:27:53 INFO mapred.MapTask: numReduceTasks: 1

22/06/15 10:27:54 INFO mapred.MapTask: (EQUATOR) 0 kvi
26214396(104857584)

22/06/15 10:27:54 INFO mapred.MapTask: mapreduce.task.io.sort.mb: 100

22/06/15 10:27:54 INFO mapred.MapTask: soft limit at 83886080

22/06/15 10:27:54 INFO mapred.MapTask: bufstart = 0; bufvoid = 104857600

22/06/15 10:27:54 INFO mapred.MapTask: kvstart = 26214396; length =
6553600

22/06/15 10:27:54 INFO mapred.MapTask: Map output collector class =
org.apache.hadoop.mapred.MapTask\$MapOutputBuffer

22/06/15 10:27:54 INFO mapred.LocalJobRunner:

22/06/15 10:27:54 INFO mapred.MapTask: Starting flush of map output

22/06/15 10:27:54 INFO mapred.MapTask: Spilling map output

22/06/15 10:27:54 INFO mapred.MapTask: bufstart = 0; bufend = 169; bufvoid
= 104857600

22/06/15 10:27:54 INFO mapred.MapTask: kvstart = 26214396(104857584);
kvend = 26214320(104857280); length = 77/6553600

22/06/15 10:27:54 INFO mapred.MapTask: Finished spill 0

22/06/15 10:27:54 INFO mapred.Task:
Task:attempt_local1115189753_0001_m_000000_0 is done. And is in the
process of committing

22/06/15 10:27:54 INFO mapred.LocalJobRunner:
hdfs://localhost:54310/rgs/test.txt:0+89

22/06/15 10:27:54 INFO mapred.Task: Task
'attempt_local1115189753_0001_m_000000_0' done.

22/06/15 10:27:54 INFO mapred.LocalJobRunner: Finishing task:
attempt_local1115189753_0001_m_000000_0

22/06/15 10:27:54 INFO mapred.LocalJobRunner: map task executor complete.

22/06/15 10:27:54 INFO mapred.LocalJobRunner: Waiting for reduce tasks

22/06/15 10:27:54 INFO mapred.LocalJobRunner: Starting task:
attempt_local1115189753_0001_r_000000_0

22/06/15 10:27:54 INFO mapred.Task: Using ResourceCalculatorProcessTree :
[]

22/06/15 10:27:54 INFO mapred.ReduceTask: Using ShuffleConsumerPlugin:
org.apache.hadoop.mapreduce.task.reduce.Shuffle@1bc68cd5

22/06/15 10:27:54 INFO reduce.MergeManagerImpl: MergerManager:
memoryLimit=334338464, maxSingleShuffleLimit=83584616,
mergeThreshold=220663392, ioSortFactor=10,
memToMemMergeOutputsThreshold=10

22/06/15 10:27:54 INFO reduce.EventFetcher:
attempt_local1115189753_0001_r_000000_0 Thread started: EventFetcher for
fetching Map Completion Events

22/06/15 10:27:54 INFO reduce.LocalFetcher: localfetcher#1 about to shuffle
output of map attempt_local1115189753_0001_m_000000_0 decomp: 211
len: 215 to MEMORY

22/06/15 10:27:54 INFO reduce.InMemoryMapOutput: Read 211 bytes from
map-output for attempt_local1115189753_0001_m_000000_0

22/06/15 10:27:54 INFO reduce.MergeManagerImpl: closeInMemoryFile ->
map-output of size: 211, inMemoryMapOutputs.size() -> 1, commitMemory ->
0, usedMemory ->211

22/06/15 10:27:54 INFO reduce.EventFetcher: EventFetcher is interrupted..
Returning

22/06/15 10:27:54 INFO mapred.LocalJobRunner: 1 / 1 copied.

22/06/15 10:27:54 INFO reduce.MergeManagerImpl: finalMerge called with 1
in-memory map-outputs and 0 on-disk map-outputs

22/06/15 10:27:54 INFO mapred.Merger: Merging 1 sorted segments

22/06/15 10:27:54 INFO mapred.Merger: Down to the last merge-pass, with 1 segments left of total size: 205 bytes

22/06/15 10:27:54 INFO reduce.MergeManagerImpl: Merged 1 segments, 211 bytes to disk to satisfy reduce memory limit

22/06/15 10:27:54 INFO reduce.MergeManagerImpl: Merging 1 files, 215 bytes from disk

22/06/15 10:27:54 INFO reduce.MergeManagerImpl: Merging 0 segments, 0 bytes from memory into reduce

22/06/15 10:27:54 INFO mapred.Merger: Merging 1 sorted segments

22/06/15 10:27:54 INFO mapred.Merger: Down to the last merge-pass, with 1 segments left of total size: 205 bytes

22/06/15 10:27:54 INFO mapred.LocalJobRunner: 1 / 1 copied.

22/06/15 10:27:54 INFO mapred.Task:

Task:attempt_local1115189753_0001_r_000000_0 is done. And is in the process of committing

22/06/15 10:27:54 INFO mapred.LocalJobRunner: 1 / 1 copied.

22/06/15 10:27:54 INFO mapred.Task: Task

attempt_local1115189753_0001_r_000000_0 is allowed to commit now

22/06/15 10:27:54 INFO output.FileOutputCommitter: Saved output of task 'attempt_local1115189753_0001_r_000000_0' to

hdfs://localhost:54310/rgs/output/_temporary/0/task_local1115189753_0001_r_000000

22/06/15 10:27:54 INFO mapred.LocalJobRunner: reduce > reduce

22/06/15 10:27:54 INFO mapred.Task: Task

'attempt_local1115189753_0001_r_000000_0' done.

22/06/15 10:27:54 INFO mapred.LocalJobRunner: Finishing task:

attempt_local1115189753_0001_r_000000_0

22/06/15 10:27:54 INFO mapred.LocalJobRunner: reduce task executor complete.

22/06/15 10:27:54 INFO mapreduce.Job: Job job_local1115189753_0001 running in uber mode : false

22/06/15 10:27:54 INFO mapreduce.Job: map 100% reduce 100%

22/06/15 10:27:54 INFO mapreduce.Job: Job job_local1115189753_0001 completed successfully

22/06/15 10:27:54 INFO mapreduce.Job: Counters: 38

File System Counters

FILE: Number of bytes read=8614

FILE: Number of bytes written=510599

FILE: Number of read operations=0

FILE: Number of large read operations=0

FILE: Number of write operations=0

HDFS: Number of bytes read=178

HDFS: Number of bytes written=69

HDFS: Number of read operations=13

HDFS: Number of large read operations=0

HDFS: Number of write operations=4

Map-Reduce Framework

Map input records=5

Map output records=20

Map output bytes=169

Map output materialized bytes=215

Input split bytes=87

Combine input records=0

Combine output records=0

Reduce input groups=10

Reduce shuffle bytes=215

Reduce input records=20

Reduce output records=10

Spilled Records=40
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=1
CPU time spent (ms)=0
Physical memory (bytes) snapshot=0
Virtual memory (bytes) snapshot=0
Total committed heap usage (bytes)=471859200

Shuffle Errors

BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters

Bytes Read=89

File Output Format Counters

Bytes Written=69

0

```
hduser@bmsce-Precision-T1700:~$ hdfs dfs -cat /1BM19CS194/output/part-00000
```

```
age 1
```

```
anitej 1
```

```
country      1
```


india	1	
is	4	
my	4	
name	1	
one	1	
prasad		1
surname		1
twenty	1	